

TiDB Analyze

A Deep Dive

Based on TiDB v8.1.0

RUSTIN LIU

Press Space to Start



Rustin Liu

PingCAP Optimizer Team Member.

Cargo/Crates.io/Rustup Maintainer.

Tokio Console Maintainer.



hi-rustin



hi-rustin.rs

Agenda

Analyze Overview

Data Structure Overview

Data Flow Overview

Data Structure & Data Flow (TiKV Perspective)

Data Structure & Data Flow (TiDB Perspective)

Q&A

Analyze Overview

Analyze Statement

-- Analyze Tables

```
ANALYZE TABLE t1, t2;
```

-- Analyze Partitions

```
ANALYZE TABLE t PARTITION p1, p2;
```

-- Analyze Columns

```
ANALYZE TABLE t COLUMNS c1, c2;
```

-- Analyze Indexes

```
ANALYZE TABLE t INDEX idx1, idx2;
```

-- Analyze Partitions' Columns

```
ANALYZE TABLE t PARTITION p1 COLUMNS c1, c2;
```

-- Analyze Partitions' Indexes

```
ANALYZE TABLE t PARTITION p1 INDEX idx1, idx2;
```

-- Analyze Predicate Columns

```
ANALYZE TABLE t PREDICATE COLUMNS;
```

-- Analyze With Only 20 Top N

```
ANALYZE TABLE t COLUMNS c1, c2 WITH 20 TOPN;
```

Data Structure Overview

Data Structure Overview

A simple example.

Create table

```
use test;  
create table t (a int);
```

Insert 2000 rows

```
import { Client } from "https://deno.land/x/mysql/mod.ts";  
  
const client = await new Client().connect({...});  
  
for (let i = 0; i < 2000; i++) {  
  await client.execute(`INSERT INTO t (a) VALUES (?)`, [i]);  
  if (i % 2 === 0) {  
    await client.execute(`INSERT INTO t (a) VALUES (?)`, [i]);  
  }  
}  
  
await client.close();
```

Data Structure Overview

Column Selectivity

```
explain select * from t where a = 100;
```

id	estRows	task	access object	operator info
TableReader_7	2.00	root		data:Selection_6
└─Selection_6	2.00	cop[tikv]		eq(test.t.a, 100)
└─TableFullScan_5	3000.00	cop[tikv]	table:t	keep order:false

```
func equalRowCountOnColumn(encodedVal []byte...) {  
    rowcount, ok := c.TopN.QueryTopN(sctx, encodedVal)  
    if ok {  
        return float64(rowcount), nil  
    }  
}
```


Data Structure Overview

Column Selectivity

TopN

```
select * from mysql.stats_top_n order by value limit 5;
```

table_id	is_index	hist_id	value	count
106	0	1	0x03800000000000000000000000000000	2
106	0	1	0x03800080000000000000000000000002	2
106	0	1	0x03800080000000000000000000000004	2
106	0	1	0x03800080000000000000000000000006	2
106	0	1	0x03800080000000000000000000000008	2

Data Structure Overview

Column Selectivity

```
explain select * from t where a = 1999;
```

id	estRows	task	access object	operator info
TableReader_7	1.00	root		data:Selection_6
└─Selection_6	1.00	cop[tikv]		eq(test.t.a, 1999)
└─TableFullScan_5	3000.00	cop[tikv]	table:t	keep order:false

```
func equalRowCountOnColumn(encodedVal []byte...) {  
    histCnt, matched := c.Histogram.EqualRowCount(sctx, val, true)  
    if matched {  
        return histCnt, nil  
    }  
}
```

Data Structure Overview

Column Selectivity

```
select hist_id, bucket_id, count, repeats,  
       CAST(lower_bound AS SIGNED) AS lower_bound,  
       CAST(upper_bound AS SIGNED) AS upper_bound,  
       ndv  
from mysql.stats_buckets order by lower_bound desc limit 5;
```

hist_id	bucket_id	count	repeats	lower_bound	upper_bound	ndv
1	229	1	1	1999	1999	0
1	228	9)	2)	1993)	1998)	0
1	227	9	2	1987	1992	0
1	226	9	2	1981	1986	0
1	225	9	2	1975	1980	0

Data Structure Overview


Histogram Bucket

- Bucket ID: The bucket ID of the histogram.
- Count: The number of values till the bucket.(**cumulative**)
- Repeats: The number of repeated values at the upper bound.
- Lower Bound: The lower bound of the bucket.
- Upper Bound: The upper bound of the bucket.
- NDV: The number of distinct values in the bucket.(**Deprecated, always 0**)

```
{  
  "bucket_id": 228,  
  "count": 9,  
  "repeats": 2,  
  "lower_bound": 1993,  
  "upper_bound": 1998,  
  "ndv": 0  
}
```

Data Structure [1]

Histogram



1. Piatetsky-Shapiro, Gregory, and Charles Connell. "Accurate Estimation Of The Number Of Tuples Satisfying A Condition."

Data Structure Overview

Column Selectivity

```
explain select * from t where a = 9999;
```

id	estRows	task	access object	operator info
TableReader_7	1.33	root		data:Selection_6
└─Selection_6	1.33	cop[tikv]		eq(test.t.a, 2000)
└─TableFullScan_5	3000.00	cop[tikv]	table:t	keep order:false

```
func equalRowCountOnColumn(encodedVal []byte...) {  
    histNDV := float64(c.Histogram.NDV - int64(c.TopN.Num()))  
    if histNDV <= 0 {  
        return 0, nil  
    }  
    return c.Histogram.NotNullCount() / histNDV, nil  
}
```

Data Structure Overview

Column Selectivity

- Not Null Count: The number of not null values in the column.
- NDV: The number of distinct values in the column.

How to calculate the NDV(Non-Distinct Value)?

We use FMSketch(Flajolet-Martin Sketch) to calculate the NDV.

Data Flow Overview

Data Flow Overview

Data Structure & Data Flow

TiKV Perspective

In TiKV, we only do two things:

1. Calculate the FMSketch.
2. Sample the data.

Data Structure - FMSketch

TiKV Perspective

Mathematical Assumptions ^[1]

1. Independence of Hash Functions:

- Assume a good hash function $h(x)$ that uniformly distributes input elements over a large range of integers.

2. Expectation of Trailing Zeros in Hash Values:

- For uniformly distributed hash values, the number of trailing zeros in their binary representation follows a geometric distribution.

1. Flajolet, Philippe; Martin, G. Nigel (1985). "Probabilistic counting algorithms for data base applications"

Data Structure - FM Sketch

TiKV Perspective

Algorithm Principles

1. Hash Mapping:

- Map each element of the set to an integer using the hash function $h(x)$.

2. Trailing Zeros Counting:

- For each hash value, count the number of trailing zeros in its binary representation. Record the maximum count R .

3. Cardinality Estimation:

- Use the maximum trailing zero count R to estimate the cardinality of the set with the formula 2^R .

Flajolet-Martin Sketch

a

b

c

d

e

f

g

h

i

j

k

l

m

n

o

p

Generate Hash Values

Flajolet-Martin Sketch - A Bad Case

a b c d e f g h i j k l m n o jj

Generate Hash Values

Data Structure - Distinct Sampling

TiKV Perspective

Core Principles ^[1]

1. Hash Function:

- Use a hash function that maps each distinct value to a random `die-level`.

2. Sample Maintenance:

- Maintain a sample S of distinct values and a current level `l`.

3. Sampling Criterion:

- Keep values in S only if their `die-level` $\geq l$.

4. Cardinality Estimation:

- Estimate distinct items as $|S| * 2^l$

1. [Phillip B. Gibbons, "Distinct Sampling for Highly-Accurate Answers to Distinct Values Queries and Event Reports"](#)

Data Structure - Distinct Sampling

TiKV Perspective

Algorithm Steps

1. Initialization:

- Start with $l = 0$ and an empty sample S .

2. Processing Each Row:

- For each row r with target attribute value v :
 - Compute die-level = $h(v)$
 - If die-level $\geq l$:
 - Add r to S

3. Sample Size Control:

- If $|S| > k$, increment l and remove items with die-level $< l$.

Distinct Sampling

Sample Size: 8

a b c d e f g h i j k l m n o jj

Process Next

Die Level: 0

Current Sample Size: 0

Estimated NDV: 0

Data Structure - Distinct Sampling

TiKV Perspective

Estimation and Analysis

1. Accuracy:

- Provides estimates within 0%-10% relative error
- Much more accurate than previous sampling methods

2. Efficiency:

- Single pass over the data.
- Only one hash function required.

Data Structure - Bernoulli Sampling

General Perspective

Mathematical Assumptions

1. Independence of Sample Selection:

- Each sample in the data set is selected independently from other samples.

2. Uniform Sampling Probability:

- Each sample is selected with a fixed probability p ($0 \leq p \leq 1$), uniformly across the entire data set.

3. Bernoulli Distribution:

- Each sample selection follows a Bernoulli distribution with parameter p .

Data Structure - Bernoulli Sampling

General Perspective

Algorithm Principles

1. Probability Definition:

- Define a sampling probability p for selecting each sample.

2. Independent Sampling:

- For each sample in the data set, generate a random number and compare it to p . If the random number is less than p , include the sample in the resulting subset.

Bernoulli Sampling

a b c d e f g h i j k l m n o p

Sample Rate: 0.3

Perform Bernoulli Sampling

Data Structure & Data Flow

TiDB Perspective

In TiDB, we do the following things:

1. Merge all FMSketches and Sample Data.
2. Build TopN and Histogram.
3. Update statistics to system tables.

Data Structure & Data Flow

Overview

Nobody can really master TiDB analyze.jpg

Configuration Name	Description	Default Value	Scope	Affected Component
<u>tidb_build_stats_concurrency</u>	The number of concurrent workers to analyze <u>tables or partitions</u>	2	Global/Session	TiDB + TiKV
<u>tidb_auto_build_stats_concurrency</u>	The number of concurrent workers to <u>automatically analyze tables or partitions</u>	1	Global (only for auto analyze)	TiDB (Owner) + TiKV
<u>tidb_analyze_distsql_scan_concurrency</u>	The number of concurrent workers to <u>scan regions</u>	4	Global/Session	TiKV
tidb_sysproc_scan_concurrency	The number of concurrent workers to scan regions	1	<u>Global (only for auto analyze)</u>	TiKV
<u>tidb_build_sampling_stats_concurrency</u>	<div>1. The number of concurrent workers to <u>merge FMSketches and Sample Data</u> from different regions</div> <div>2. The number of concurrent workers to <u>build TopN and Histogram</u></div>	2	Global/Session	TiDB
<u>tidb_analyze_partition_concurrency</u>	The number of concurrent workers to <u>save statistics to the system tables</u>	2	Global/Session	TiDB

Data Structure & Data Flow

TiDB Perspective - Analyze tables or partitions concurrently

`tidb_build_stats_concurrency`

Data Structure & Data Flow

TiDB Perspective - Scan regions concurrently

`tidb_analyze_distsql_scan_concurrency`

Data Structure & Data Flow

TiDB Perspective - Merge FMSketches and Sample Data

`tidb_build_sampling_stats_concurrency`

Data Structure & Data Flow

TiDB Perspective - Build TopN and Histogram

`tidb_build_sampling_stats_concurrency`

Build TopN

Count: 200 Sample Length: 23

Samples

TopN List

104	23	45	164	57	34	212	92
69	176	116	140	116	128	80	152
188	224	13	200	200	5	5	

Current:

Current Count: 0


Step:

Next Step

Reset

Build Histogram

Count: 200 Sample Length: 17



Sample Factor: 11.8 Per Bucket: 78.4 Current Sample: N/A Current Bucket: 1

Step: Click Next Step to start

Next Step

Reset

Data Structure & Data Flow

TiDB Perspective - Save Statistics

`tidb_analyze_partition_concurrency`

Data Structure & Data Flow

TiDB Perspective - Merge Global Statistics

`tidb_merge_partition_stats_concurrency`

Improve the Analyze

What can we do?

Tracking Document

Statistics Project Planning and Implementation

Statistics Tech Debt

Blogs

I started a new series blog post to discuss the Analyze feature improvement.

NCRMTA1: Surprise analyze-partition-concurrency-quota

NCRMTA2: Accelerate Auto-Analyze of Partitioned Tables

Q&A

Do you have any questions?

Thank You!